

Performance comparison: a table lists execution time, speedup (using mode 0 as the baseline) of each mode the excludes the data initialization and result verification.

I chose to implement merge sort.

| Mode | Execution Time | Ratio to the base(mode 0) |
|------|----------------|---------------------------|
| 0 | 27046002 | 1 |
| 1 | 20271504 | 0.75 |
| 2 | 28651328 | 1.05 |

What's the complexity of your algorithm?

Complexity of merge sort in openCL is $O(\log(n) * \log(\text{stride}))$. OpenCL sorts chunks (say, 256) and then does the merging. The entire thing is done for $\log(n)$ times.

Is the performance improvement as good as you expected? Why or why not?

Ideally, the algorithm should provide much better results. It should be improved by a factor of 2 at least. But my code uses local memory, so synchronization is an overhead and it degrades the performance of parallel sorting.

Comparing your performance with pthread

| Name | Execution Time | Ratio to the base(pthread) |
|-----------------------|----------------|----------------------------|
| Pthread merge sorting | 20271504 | 1 |
| FPGA merge sorting | 28651328 | 1.4 |

Performance should be better in openCL than compared to pthread. Because for 2 cores, pthread has two threads for each of the core. And, FPGA/openCL can have a large number of work units to work on a large number of elements in a parallel environment. Here, we are getting little slower performance because of overhead.

The list of students sharing the same board with you :

- Sweta Rout (srout@ncsu.edu)
- Cody Nesbitt (cznesbit@ncsu.edu)